

Modular implicits / Multi-core OCaml

Leo White

Modular implicits

Leo White, Frédéric Bour and Jeremy Yallop

Ad-hoc polymorphism

Ad-hoc polymorphism occurs when a function is defined over several different types, acting in a different way for each type.

`4 + 9`

`4.5 + 9.5`

`print [true; false]`

`print (Some 8.4)`

Simple Overloading

```
public static String show(int x)
{
    return Integer.toString(x);
}
```

```
public static String show(float x)
{
    return Float.toString(x);
}
```

Simple Overloading

```
show (7);
```

```
show (4.2);
```

```
show("foo");
```

```
error: no suitable method found for show(String)
      show("foo")
      ^
```

Simple Overloading

```
public static <T> String show_twice(T x)
{
    return show(x) + " " + show(x);
}
```

```
Main.java:23: error: no suitable method found for show(T)
        return show(x) + " " + show(x);
               ^
```

method Show.show(int) is not applicable

(argument mismatch; T cannot be converted to int)

method Show.show(float) is not applicable

(argument mismatch; T cannot be converted to float)

where T is a type-variable:

T extends Object declared in method <T>show_twice(T)

Type classes

```
class Show a where  
  show :: a -> string
```

```
instance Show Int where  
  show = showInt
```

```
instance Show Float where  
  show = showFloat
```

Type classes

```
> show 7  
"7"
```

```
> show 4.5  
"4.5"
```

```
> show (\ x -> x)
```

No instance for (Show (t1 -> t1)) arising
from a use of 'show'

Type classes

```
show_twice x = show x ++ " " ++ show x
```

Type classes

```
show_twice :: Show a => a -> String
```

```
show_twice x = show x ++ " " ++ show x
```

Type classes

```
> show_twice 5
```

```
"5 5"
```

```
> show_twice (\ x -> x)
```

```
No instance for (Show (t0 -> t0))  
  arising from a use of 'show_twice'
```

Type classes

```
instance Show a => Show [a] where  
  show = showList show
```

```
> show [7, 8, 9]  
" [7,8,9]"
```

```
> show [[1, 2, 3], [4, 5, 6]]  
" [[1,2,3],[4,5,6]]"
```

Coherence

Every different valid typing derivation for a program leads to a resulting program that has the same dynamic semantics.

Canonicity

```
> instance Show Int where  
    show x = "An Int"
```

```
<interactive>:2:10:
```

Duplicate instance declarations:

```
instance Show Int -- Defined at <interactive>:2:10  
instance Show Int -- Defined in 'GHC.Show'
```

Abstract type equalities

```
module M : sig
  type t
end = struct
  type t = int
end
```

Abstract type equalities

```
module F (X : sig type t val show : t -> string end) =  
  struct  
    instance Show X.t where  
      show = X.show  
  end
```

```
instance Show int where  
  show = string_of_int
```

```
F(struct  
  type t = int  
  let show _ = "An int"  
end)
```


Scala implicits

```
trait Showable[T] { def show(x: T): String }

def show[T](x: T)(implicit s: Showable[T]) =
  s.show(x)

implicit object IntShowable extends Showable[Int] {
  def show(x: Int) = x.toString
}

show(7)
```

Scala implicits

```
def show_twice[T](x: T)(implicit s: Showable[T]) =  
  show(x) + " " + show(x)
```

```
show_twice(7)
```

Scala implicits

```
implicit class ListShowable[T]  
  extends Showable[List[T]]  
    (implicit s: Showable[T]) {  
  
  def show(x: List[T]) = x.toString(s.show, x)  
  
}  
  
show(List(1,2,3))
```

Scala implicits

```
implicit object IntShowable2 extends Showable[Int] {  
  def show(x: Int) = x.toString  
}
```

show(7)

error: ambiguous implicit values:

both object IntShowable2 in object \$iw of type
object IntShowable2

and object IntShowable in object \$iw of type
object IntShowable

match expected type Showable[Int]
show(7)

Modular implicits

Implicit *module* parameters to functions chosen by their *module type*.

Demo

Status

Working prototype based on OCaml 4.02.0 (by Frédéric Bour)

- ▶ Install it using the OCaml Package Manager (OPAM):
`$ opam switch 4.02.0+modular-implicits`
- ▶ Try it online (all compiled to JavaScript and running in the browser):
<http://andrewray.github.io/iocamljs/modimp.html>
- ▶ When you (inevitably) find bugs, report them to
<http://github.com/ocaml-labs/ocaml-modular-implicits>

Multi-core OCaml

Stephen Dolan, Leo White, KC Sivaramakrishnan and Anil Madhavapeddy

Concurrency and Parallelism

Concurrency

Parallelism

Concurrency and Parallelism

Concurrency

- ▶ Concurrency is for writing programs
“My program handles 1000s connections at once”

Parallelism

Concurrency and Parallelism

Concurrency

- ▶ Concurrency is for writing programs
“My program handles 1000s connections at once”

Parallelism

- ▶ Parallelism is for improving performance
“My program uses all 8 cores”

Concurrency and Parallelism

Concurrency

- ▶ Concurrency is for writing programs
“My program handles 1000s connections at once”
- ▶ Monadic: Lwt, Async

Parallelism

- ▶ Parallelism is for improving performance
“My program uses all 8 cores”

Concurrency and Parallelism

Concurrency

- ▶ Concurrency is for writing programs
“My program handles 1000s connections at once”
- ▶ Monadic: Lwt, Async
- ▶ Direct: Systhreads, Vmthreads

Parallelism

- ▶ Parallelism is for improving performance
“My program uses all 8 cores”

Concurrency and Parallelism

Concurrency

- ▶ Concurrency is for writing programs
“My program handles 1000s connections at once”
- ▶ Monadic: Lwt, Async
- ▶ Direct: Systhreads, Vmthreads

Parallelism

- ▶ Parallelism is for improving performance
“My program uses all 8 cores”
- ▶ Multi-process: Parmap, Async_parallel

Concurrency and Parallelism

Concurrency

- ▶ Concurrency is for writing programs
“My program handles 1000s connections at once”
- ▶ Monadic: Lwt, Async
- ▶ Direct: Systhreads, Vmthreads

Parallelism

- ▶ Parallelism is for improving performance
“My program uses all 8 cores”
- ▶ Multi-process: Parmap, Async_parallel
- ▶ Shared memory: ?

Multi-core OCaml


- ▶ Provide support for shared-memory parallelism
- ▶ Improve support for concurrency – avoid people abusing the parallelism primitives for concurrency (see Java).

Parallelism: Domains

The unit of parallelism is the *domain*.

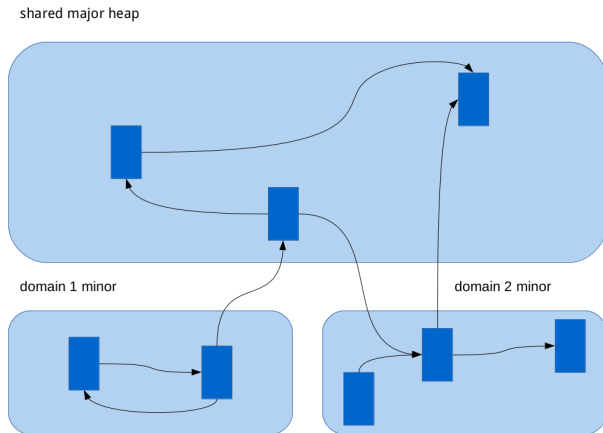
- ▶ Expensive to create
- ▶ Intention is to have roughly one per-core

Minor heaps




Each domain has its own minor heap

Minor heaps



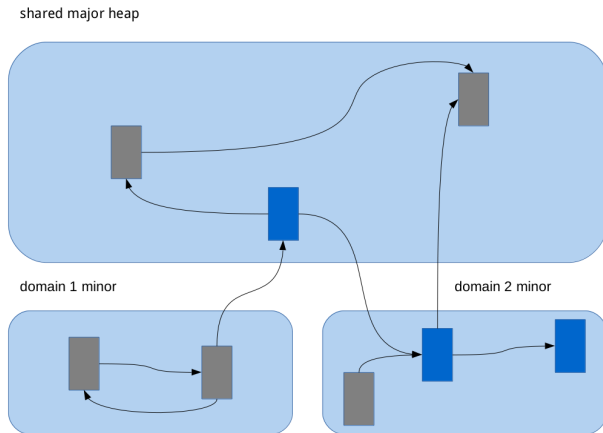
These minor heaps can be collected independently without synchronisation

Minor heaps



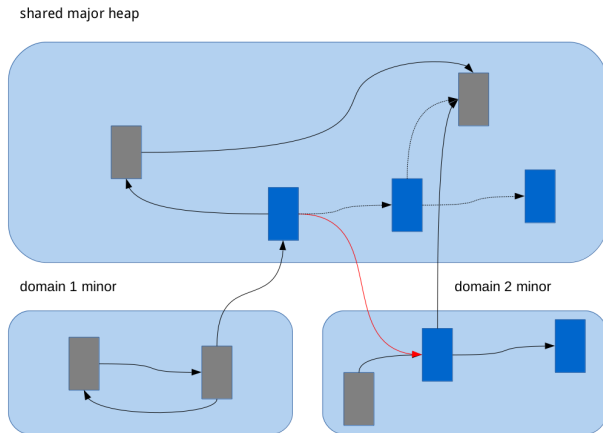
GC invariant: no pointers between minor heaps

Minor heaps



GC invariant: no pointers between minor heaps

Minor heaps



GC invariant: no pointers between minor heaps

Major heap

Mostly-concurrent parallel collector (VCGC)

- ▶ Domains independently mark reachable objects




- ▶ Domains sweep separate parts of the heap



Major heap

Complete a GC cycle by changing the interpretation of the mark bits



- ▶ Requires all domains to synchronise
- ▶ Most marking and sweeping should have been completed before synchronisation

Concurrency: Fibers

The unit of concurrency is the *fiber*.

- ▶ Very cheap to create
- ▶ A fiber is essentially just a stack
- ▶ Stacks start very small and are automatically resized as needed

Concurrency: Fibers

There are many interesting programming models for concurrency.

- ▶ We don't want to mandate a particular model
- ▶ Instead provide powerful primitives for implementing concurrency

Concurrency: Fibers

There are many interesting programming models for concurrency.

- ▶ We don't want to mandate a particular model
- ▶ Instead provide powerful primitives for implementing concurrency
- ▶ *Algebraic effects*

Demo

Status

Prototype based on OCaml 4.02.1 (by Stephen Dolan)

- ▶ Install it using the OCaml Package Manager (OPAM):

```
$ opam remote add ocamlabs git://github.com/ocamlabs/opam-  
$ opam switch 4.02.1+multicore
```

- ▶ Byte-code only at the moment
- ▶ GC needs testing, tuning and benchmarking
- ▶ Some features broken (weak references, finalizers, lazy values)
- ▶ When you (inevitably) find bugs, report them to
<http://github.com/ocamlabs/ocaml-multicore>